

# Practice activity: Preparing a model for deployment

 [coursera.org/learn-foundations-of-ai-and-machine-learning-supplement/ChkxN/practice-activity-preparing-a-model-for-deployment](https://www.coursera.org/learn-foundations-of-ai-and-machine-learning-supplement-ChkxN/practice-activity-preparing-a-model-for-deployment)

## Introduction

In this activity, you will apply the skills and knowledge you've gained to prepare an ML model for deployment. This exercise will give you hands-on experience with the steps involved in packaging, containerizing, and deploying a model in a production environment.

By the end of this lab, you will be able to:

- Take a trained model, prepare it for deployment, and ensure it's ready to serve predictions in a real-world setting.

## Objective

- Package an ML model along with its dependencies.
- Containerize the model using Docker.
- Deploy the model locally, and test its performance.
- Optionally, deploy the model to a cloud environment for wider accessibility.

## Prerequisites

Before starting the activity, ensure you have the following tools installed on your machine:

- **Python 3.6 or later:** for running scripts and packaging the model
- **Docker:** for containerizing the model
- **Git:** (optional) for version control

- An **IDE**: such as VS Code, PyCharm, or Jupyter Notebook

## 1. Packaging the model

---

### Step-by-step guide:

---

#### Step 1: Save the trained model

---

##### 1. Train a simple model

If you don't have a trained model, you can use the following example to train a simple logistic regression model on the Iris dataset:

```
1
2
3
4
5
6
7
8
9
10
11
12
13
14
15
from sklearn.datasets import load_iris
from sklearn.model_selection import train_test_split
```

```
from sklearn.linear_model import LogisticRegression  
import joblib  
  
# Load the dataset  
  
iris = load_iris()  
  
X_train, X_test, y_train, y_test = train_test_split(iris.data, iris.target, test_size=0.2, random_state=42)  
  
# Train the model  
  
model = LogisticRegression(max_iter=200)  
  
model.fit(X_train, y_train)  
  
# Save the model to a file  
  
joblib.dump(model, 'iris_model.pkl')
```



## 2. Ensure the model is saved

Verify that the model file iris\_model.pkl is saved in your working directory.

## Step 2: Create a requirements.txt file

---

### 1. List the dependencies

Create a requirements.txt file to list all the Python packages your model depends on. This ensures that the same environment can be recreated in production.

Example requirements.txt:

4

flask==2.0.2



### 2. Save the file

Save the requirements.txt file in the same directory as your model.

## Step 3: Create a Python script for serving the model

---

### 1. Create a serve\_model.py script

This script will use Flask to serve the model as a RESTful API, allowing it to make predictions when called.

1

2

3

4

5

6

7

8

9

10

11

12

13

14

15

16

17

```
from flask import Flask, request, jsonify
```

```
import joblib
```

```
import numpy as np
```

```
app = Flask(__name__)
```

```
# Load the model
```

```
model = joblib.load('iris_model.pkl')

@app.route('/predict', methods=['POST'])

def predict():

    data = request.get_json(force=True)

    prediction = model.predict(np.array(data['input']).reshape(1, -1))

    return jsonify({'prediction': int(prediction[0])})

if __name__ == '__main__':

    app.run(host='0.0.0.0', port=80)
```



## 2. Test the script

Run the script using `python serve_model.py`, and use tools such as Postman or cURL to test the API locally by sending a POST request to `http://127.0.0.1:80/predict` with a JSON body, such as:

```
1  
2  
3  
{  
    "Input" : [5.1, 3.5, 1.4, 0.2]  
}
```



## 2. Containerizing the model with Docker

---

### Step-by-step guide

---

#### Step 1: Create a Dockerfile

---

##### 1. Write a Dockerfile

The Dockerfile defines the environment in which your model will run, including the base image, dependencies, and startup command.

Example Dockerfile:

1

2

3

4

5

6

7

8

9

10

11

12

13

14

15

16

17

```
# Use an official Python runtime as a parent image
```

```
FROM python:3.8-slim
```

```
# Set the working directory in the container
```

```
WORKDIR /app
```

```
# Copy the current directory contents into the container at /app
```

```
COPY . /app

# Install any needed packages specified in requirements.txt

RUN pip install --no-cache-dir -r requirements.txt

# Make port 80 available to the world outside this container

EXPOSE 80

# Run serve_model.py when the container launches

CMD [python, serve_model.py]
```



## 2. Save the Dockerfile

Save it in the same directory as your model and scripts.

### **Step 2: Build the Docker image**

#### **1. Build the image**

Open a terminal, navigate to your project directory, and build the Docker image using the following command:

1

```
docker build -t iris_model_image .
```



## 2. Verify the image

Once the build is complete, verify that the image was created by running docker images.

## Step 3: Run the Docker container

---

### 1. Run the container

Run the Docker container from the image you just created:

1

```
docker run -d -p 80:80 iris_model_image
```



## 2. Test the deployment

Test the deployed model by sending a POST request to `http://localhost:80/predict` using the same JSON payload you used earlier.

## 3. Optional: Deploying to a cloud environment

---

If you want to deploy your Docker container to a cloud environment (such as Microsoft Azure, Amazon Web Services, or Google Cloud), follow the platform-specific instructions for deploying Docker containers. Here's a brief overview for deploying to Azure:

### Step-by-step guide:

---

#### Step 1: Push the Docker image to Azure Container Registry

---

##### 1. Log in to Azure

Log in to your Azure account using the Azure command-line interface:

##### 2. Create a resource group and container registry

Create a resource group and container registry to store your Docker image:

2

```
az group create --name myResourceGroup --location eastus
```

```
az acr create --resource-group myResourceGroup --name myContainerRegistry --  
sku Basic
```



### 3. Push the image

Tag your Docker image, and push it to the Azure Container Registry:

1

2

```
docker tag iris_model_image myContainerRegistry.azurecr.io/iris_model_image:v1
```

```
docker push myContainerRegistry.azurecr.io/iris_model_image:v1
```



## Step 2: Deploy the container to Azure Container instances

---

### 1. Deploy the container

- Deploy the container directly from your Azure Container Registry:

1

```
az container create --resource-group myResourceGroup --name irisModelContainer --  
image myContainerRegistry.azurecr.io/iris_model_image:v1 --dns-name-  
label irismodel --ports 80
```

## 2. Test the deployed model

Once deployed, you can test the model using the public IP address or DNS label provided by Azure.

## 4. Deliverables

---

Upon completing this activity, you should have your packaged model, Dockerfile, and any scripts you used. If you deployed the model to the cloud, provide the URL to access it.

**Reflection:** Write a brief reflection on the challenges you encountered and how you overcame them during the lab.

## Conclusion

---

This activity provides you with hands-on experience in preparing an ML model for deployment. By packaging, containerizing, and deploying the model, you gain a deeper understanding of the steps required to transition from model development to production. These skills are crucial for ensuring that your models can be effectively deployed and scaled in real-world environments.



**Hi, Rodrigo!**

---

**How can I help?**

---